

Assessed Exercise 1 (Mandatory): Parsing, Code Generation and State Manipulation in Haskell: a Real-world Application

Functional Programming 4

School of Computing Science, University of Glasgow

Hand-out/hand-in: see Moodle

The problem

The objective of this coursework is to assess your understanding of functional programming and in particular of the concepts of higher-order functions, list manipulation, state and input/output, as well as the use of libraries, code organisation and testing.

The aim of this coursework is to write a program in Haskell which will analyse a text file for specific patterns and produce a new text file according to a number of templates. To be specific, your program will analyse Fortran-95 code which has been annotated with OpenACC pragmas, and generate code for GPU acceleration using the OpenCL API. This is a real-world application of parsing, code generation and state manipulation.

However, you do not need to know what OpenACC or OpenCL are to do this coursework. You will be provided with a code skeleton and an example input file. Your parser should be able to process not just the example input file but similar files as well, as your final code will be tested against a similar but different file.

What to hand in

Your solution should be an archive (.tgz, .tbz2 or .zip) containing a set of Haskell source code files based on the provided code skeleton.

Include full identification information at the beginning of the file in a block comment(`{- -}`), one entry per line: your name and matriculation number (one line), the course (Functional Programming 4), the exercise title, and the date. These should all be inside a comment.

After the identification, include a short status report, also inside a comment. This isn't documentation of the program; it simply *reports the status of your submission*. Which parts did you do? Does your solution work? What bugs are you aware of? Which extensions, if any, did you implement?

After the status report comes the program.

Although it is essential that you perform unit tests on your code, you don't have to submit the unit test code.

If your code can't be built with the command `ghc --make Main` then you must provide a Cabal build file.

Submit your code archive on Moodle, before the hand-in deadline.

Assessment

The entire exercise is worth 20% of the course assessment, or 15% if you also do the optional second coursework, and it will be marked in bands.

The assessment is *entirely* based on unit tests. For each function provided in the code skeleton, one or more unit tests will be run to evaluate the correctness of the implementation. The mark you get is based solely on the number of unit tests passed.

Code skeleton and example source files

The code is contained in a number of Haskell source files, in the archive `Skeleton.zip` on Moodle. Each of the files contains a number of functions and/or datatypes. For each function the type declarations and a stub are provided.

The archive `Fortran95-sources.zip` contains the Fortran 95 source code template `module_LES_ocl_TEMPL.f95` and an example of the generated code, `module_LES_ocl.f95`. You should study both source files to understand what information you need to extract from the template and what code you need to generate.

Some more information on Fortran-95, OpenACC and OpenCL

I will use simplified regular expressions below to describe the format of the code. I use `{regex}` to enclose the regular expressions. So, for example

- `{long|short}` means a choice between 'long' and 'short';
- `{[1-8]D}` means a choice of a digit from 1 to 8, followed by 'D';
- `{unsigned?}` means one or zero occurrences of 'unsigned';
- `{[0-9]+}` means one or more digits between 0 and 9, i.e. an unsigned integer value
- `${var}` is used to indicate the name of a variable to be substituted

Fortran-95 syntax

- Comments start with '!', there are no block comments
- Whitespace is not significant.
- Fortran is NOT case sensitive.
- Statements must fit on a single line, but lines can be broken up using '&'.

- If you want to put multiple statements on a single line, they must be separated by semicolons; otherwise no semicolons.
- Apart from the actual type, the other attributes of a type declaration are optional; their order is arbitrary. However, you can assume the order as used in the example code.

OpenACC pragmas

The Fortran template code contains OpenACC pragmas of two forms:

1. Pragmas for code analysis

```
!$ACC {Const?}Arguments
    <variable declaration>
    ....
!$ACC End {Const?}Arguments
```

These define regions of code (in the case of the coursework, variable declarations) that are required by the code generator. Some examples of valid variable declarations:

```
real(kind=4), dimension(-1:ip+1,-1:jp+1,0:kp+1) :: sm1, sm2, sm3
real(kind=4), dimension(kp+2) :: z2
real(kind=4) :: dt
integer :: im
real(kind=8), intent(inout) :: sor_err
```

Moreover, each variable declaration can *optionally* be followed by an OpenACC pragma

```
!$ACC Argmode {Read|Write|ReadWrite}
```

For example

```
integer(kind=8), dimension(0:ip+1,0:jp+1,0:kp+1), intent(in) :: t1,t2,t3 !$ACC A
```

2. Pragmas for code generation

```
!$ACC {BufDecls|SizeDecls|MakeSizes|MakeBuffers|SetArgs|WriteBuffers}
```

These mark the place where the generated code should be inserted and identify what type of code should be generated.

Code generation and OclWrapper OpenCL API

For each of the keywords in the pragmas for code generation above, specific code has to be generated. We use a simple naming convention: for a variable name $\{var\}$, the corresponding buffer will be name $\{var\}_{buf}$ and the corresponding size $\{var\}_{sz}$.

- BufDecls

Buffer declarations are of the form

```
integer(8) ::  $\{var\}_{buf}$ 
```

- SizeDecls

Size declarations are of the form

```
integer, dimension( $\{1-4\}$ )::  $\{var\}_{sz}$ 
```

- MakeSizes

There are two ways to create the sizes: generate-time computed and run-time computed.

- For run-time computed, the form is

```
 $\{var\}_{sz} = \text{shape}(\{var\})$ 
```

For generate-time computed, the form is a Fortran array constant:

```
 $\{var\}_{sz} = (/ \{([0-9]+, )*[0-9]+\} /)$ 
```

where the size has been pre-computed by the code generator, for example

```
p_sz = (/ 150, 150, 90 /)
```

This means the code generator has to interpret the expressions for the dimensions of the arrays.

For generating the OpenCL code we use OclWrapper¹, an API to facilitate integration of OpenCL in large Fortran-95 codebases.

For the purpose of the coursework, the following API calls are required:

- MakeBuffers

```
call oclMake $\{([1-4])D\}\{Int/Float\}Array\{Read/Write/ReadWrite\}Buffer(\{var\}_{buf}, \{var\}_{sz})$ 
```

- SetArgs

```
call oclSet $\{Int/Float\}\{Const/Array\}Arg(\{([0-9]+), \{var\}_{buf} \}$ 
```

Here, the integer constant reflects the order of the variables in the source code, i.e. the first encountered variable has argument position 0, the second one has position 1 and so on.

¹<https://github.com/wimvanderbauwhede/OpenCLIntegration>

- WriteBuffers

```
call oclWrite{[1-4]D}{Int/Float}ArrayBuffer(${var}_buf, ${var}_sz, ${var})
```